

# Deep Learning Software

## PyTorch

Department of Computer Science, NYCU

TA 劉子齊 Jonathan

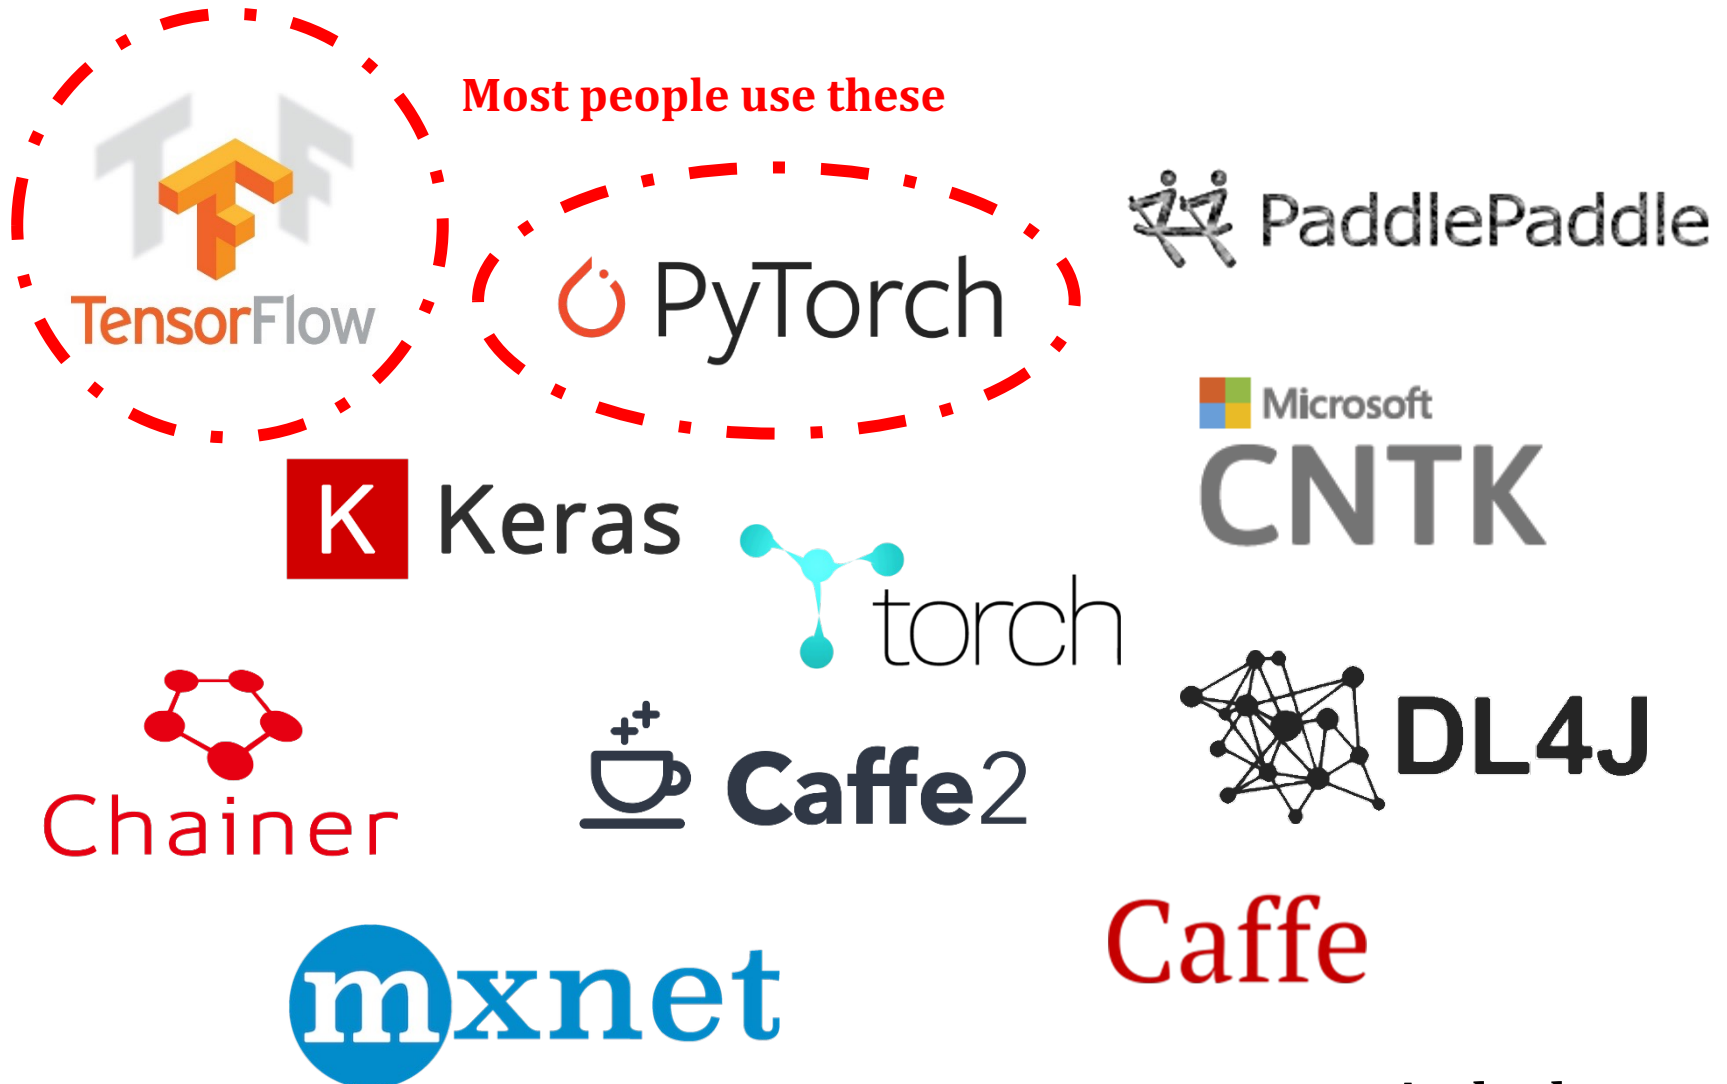
Reference: Stanford CS231n

# Frameworks



Caffe

# Frameworks



# Frameworks

We will focus on this



Caffe

# Advantages of PyTorch

- Developing and testing new ideas are quickly
- Computing gradients automatically
- Running model structures on GPU is efficiently

**Please use **PyTorch** in this course !!**

 PyTorch  PyTorch  PyTorch

# Computational Graphs


## Numpy

```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)
```



Neural network can be denoted as a directed acyclic graph

# Computational Graphs

## Numpy

```
import numpy as np
np.random.seed(0)


N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
```

```
a = x * y
b = a + z
c = np.sum(b)
```

```
grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```

compute gradients



## Problems:

- Can't run on GPU
- Have to compute our own gradients

# Computational Graphs

## Numpy


```
import numpy as np
np.random.seed(0)
```

```
N, D = 3, 4
```

```
x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)
```

```
a = x * y
b = a + z
c = np.sum(b)
```

```
grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



## PyTorch

```
import torch
```

```
N, D = 3, 4
```

```
x = torch.randn(N, D)
y = torch.randn(N, D)
z = torch.randn(N, D)
```

```
a = x * y
b = a + z
c = torch.sum(b)
```

Looks exactly like numpy!



# Computational Graphs

## Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



## PyTorch

```
import torch

N, D = 3, 4
x = torch.randn(N, D, requires_grad=True)
y = torch.randn(N, D)
z = torch.randn(N, D)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

PyTorch handles gradients for us!

.backward() computes the gradient

# Computational Graphs

## Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



## PyTorch

```
import torch
device = 'cuda:0'
N, D = 3, 4
x = torch.randn(N, D, requires_grad=True, device=device)
y = torch.randn(N, D, device=device)
z = torch.randn(N, D, device=device)

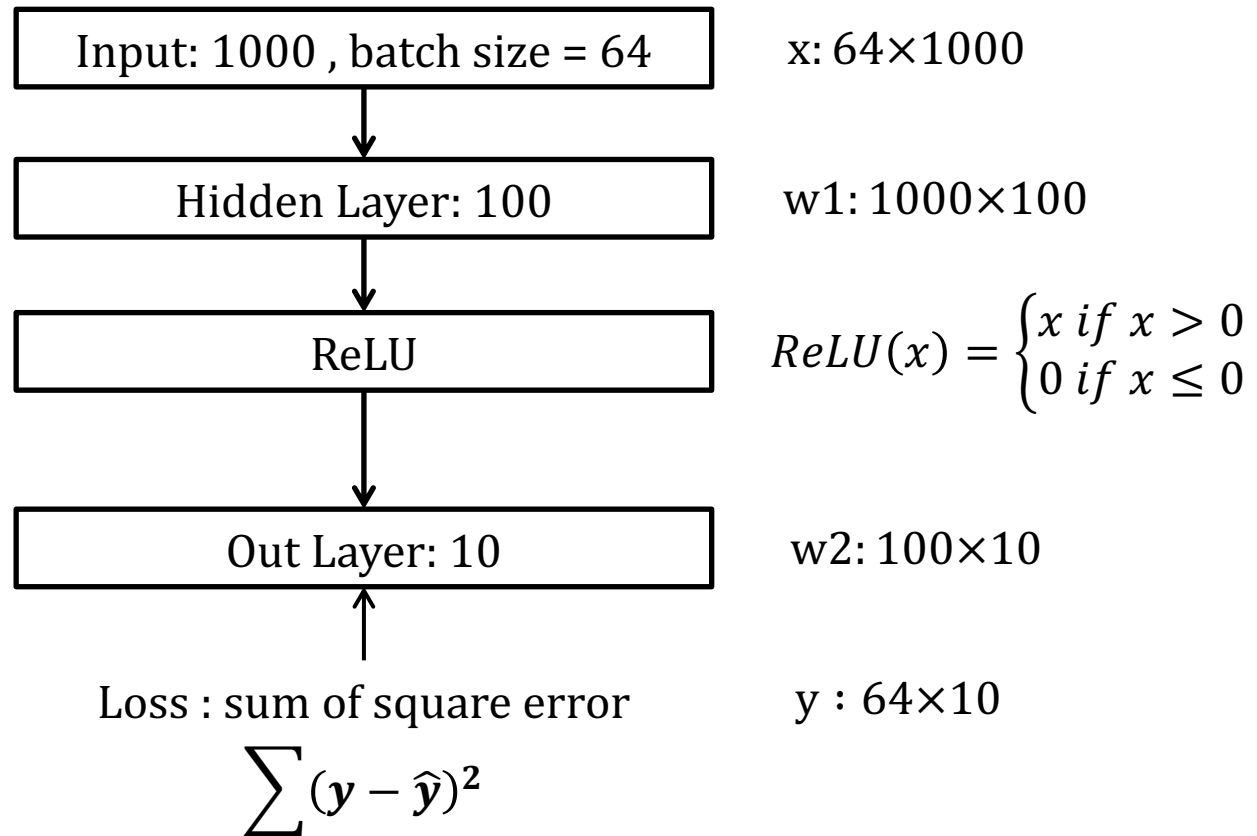
a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

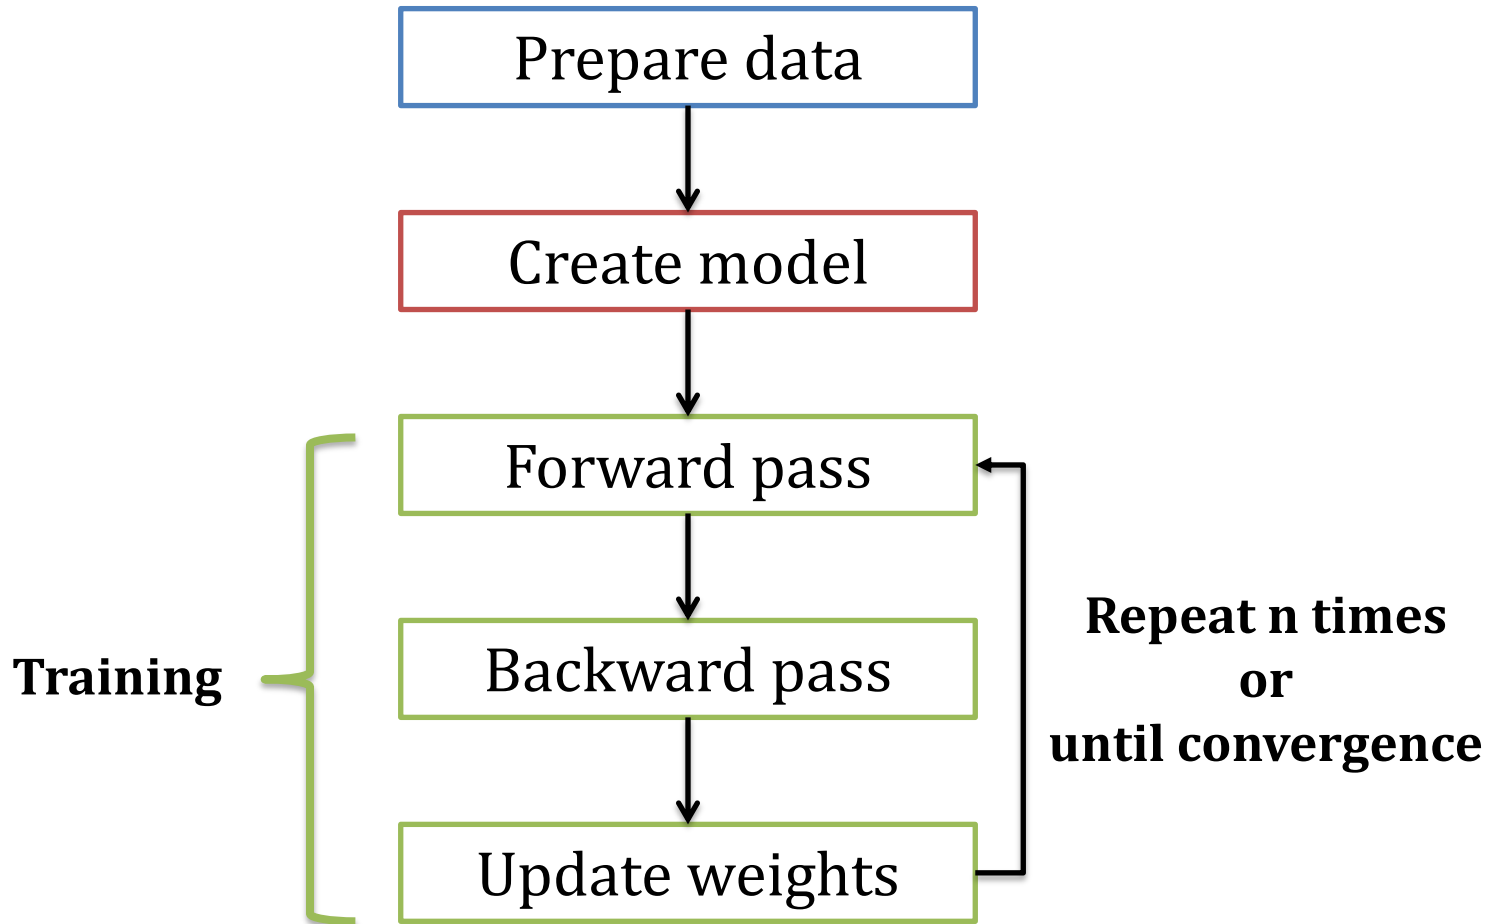
Trivial to run on GPU - just construct arrays on a different device!

# Example

- 2-layer network



# Flow Chart



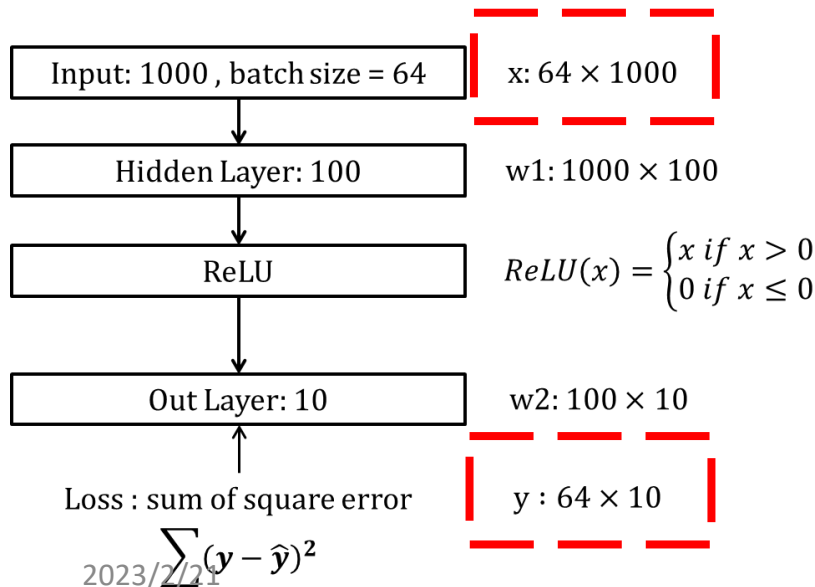
# Step1. Prepare Data

## PyTorch Tensors

Create random tensors as input and ground truth

To run on GPU, just use a different device, like a following:

```
device = torch.device('cuda:0')
```



```
import torch
```

```
device = torch.device('cpu')
learning_rate = 1e-6
```

```
x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
```

```
w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)
```

```
for t in range(300):
```

```
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y)
```

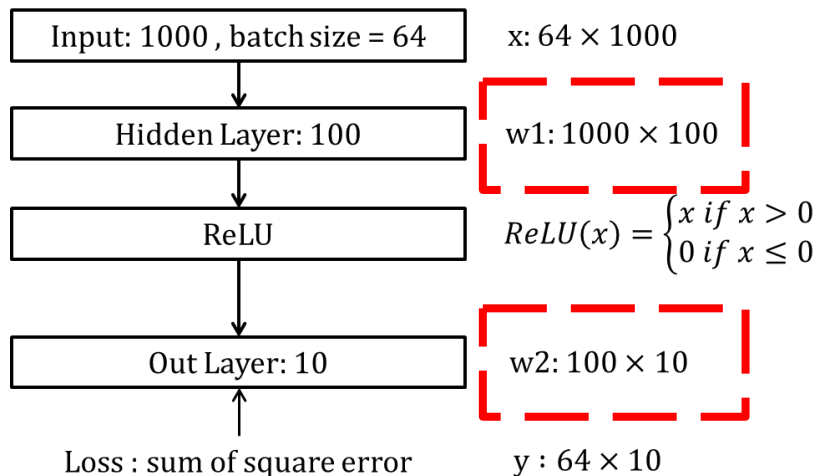
```
    grad_y_pred = 2.0 * loss
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)
```

```
    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```

```
print(loss.pow(2).sum())
```

## Step2. Create Model PyTorch Tensors

Create random tensors as layer weights



$$\sum (y - \hat{y})^2$$

```
import torch
```

```
device = torch.device('cpu')  
learning_rate = 1e-6
```

```
x = torch.randn(64, 1000, device=device)  
y = torch.randn(64, 10, device=device)
```

```
w1 = torch.randn(1000, 100, device=device)  
w2 = torch.randn(100, 10, device=device)
```

```
for t in range(300):
```

```
    h = x.mm(w1)  
    h_relu = h.clamp(min=0)  
    y_pred = h_relu.mm(w2)  
    loss = (y_pred - y)
```

```
    grad_y_pred = 2.0 * loss  
    grad_w2 = h_relu.t().mm(grad_y_pred)  
    grad_h_relu = grad_y_pred.mm(w2.t())  
    grad_h = grad_h_relu.clone()  
    grad_h[h < 0] = 0  
    grad_w1 = x.t().mm(grad_h)
```

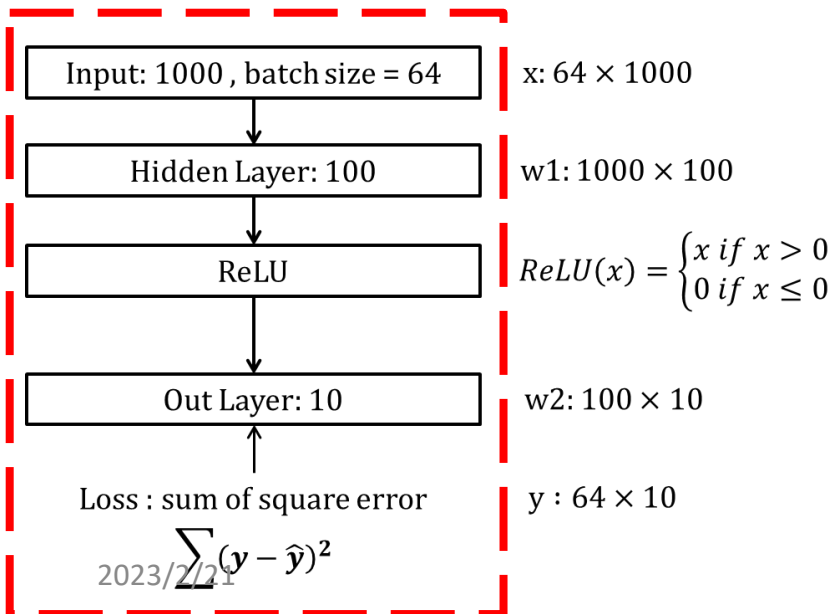
```
    w1 -= learning_rate * grad_w1  
    w2 -= learning_rate * grad_w2
```

```
print(loss.pow(2).sum())
```

# Step3. Forward pass

## PyTorch Tensors

Compute predictions and loss



```
import torch

device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for t in range(300):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * loss
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2


    print(loss.pow(2).sum())
```



# Step4. Backward pass

## PyTorch Tensors

Manually compute gradients



```
import torch

device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for t in range(300):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y)

    grad_y_pred = 2.0 * loss
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2


    print(loss.pow(2).sum())
```



# Step5. Update Weights

## PyTorch Tensors

Gradient descent step on weights



```
import torch

device = torch.device('cpu')
learning_rate = 1e-6

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)

w1 = torch.randn(1000, 100, device=device)
w2 = torch.randn(100, 10, device=device)

for t in range(300):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y)

    grad_y_pred = 2.0 * loss
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2

print(loss.pow(2).sum())
```

# Implement your own DL model by using **PyTorch**

# Step1. Prepare Data

## PyTorch.utils.data

**DataLoader** wraps a **Dataset** and provides minibatches, shuffling, multithreading, for you

When you need to load custom data, just write your own Dataset class

Iterate over loader to form minibatches

<https://github.com/utkuozbulak/pytorch-custom-dataset-examples>

2023/2/21

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred, y_batch)
        print(loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()
```

## Step2. Create Model

### PyTorch.nn

Higher-level wrapper for working with neural nets

Use this ! It will make your life easier

A PyTorch Module is a neural net layer, it can contain weights or other modules

Define your whole model as a single module

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred, y_batch)
        print(loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()
```

## Step2. Create Model

### PyTorch.nn

Initializer sets up two children  
(Module can contain Modules)

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(),
                               lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)
        print(loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()
```

## Step2. Create Model

### PyTorch.nn

Define forward pass using child modules

No need to define backward – autograd will handle it

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(),
                               lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred,
                                              y_batch)
        print(loss.item())

        loss.backward()

    optimizer.step()
    optimizer.zero_grad()
```

## Step3. Forward pass

### PyTorch.nn

Define forward pass using child modules

Feed data to model, and compute loss

nn.functional has useful helpers like loss functions

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred, y_batch)
        print(loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()
```

## Step4. Backward pass

### PyTorch.autograd

Forward pass looks exactly the same as before, but we don't need to track intermediate values

PyTorch keeps track of them for us in the computational graph

Compute gradient of loss with respect to all model weights (they have `requires_grad=True`)

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred, y_batch)

        print(loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()
```



# Step5. Update Weights

## PyTorch.optim

Use an **optimizer** for different update rules

After computing gradients, use optimizer to update each model parameters and reset gradients

```
import torch
from torch.utils.data import TensorDataset, DataLoader

device = torch.device('cpu')
learning_rate = 1e-2

x = torch.randn(64, 1000, device=device)
y = torch.randn(64, 10, device=device)
loader = DataLoader(TensorDataset(x, y), batch_size=8)

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear_1 = torch.nn.Linear(D_in, H)
        self.linear_2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h = self.linear_1(x)
        h_relu = torch.nn.functional.relu(h)
        y_pred = self.linear_2(h_relu)
        return y_pred

model = TwoLayerNet(D_in=1000, H=100, D_out=10)
model = model.to(device)

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

for epochs in range(50):
    for x_batch, y_batch in loader:
        y_pred = model(x_batch)
        loss = torch.nn.functional.mse_loss(y_pred, y_batch)

        print(loss.item())

        loss.backward()

        optimizer.step()
        optimizer.zero_grad()
```

# Real Application

- MNIST example for PyTorch



- <https://github.com/pytorch/examples/tree/master/mnist>

# Build and train a CNN classifier

- Data Loader
- Define Network
- Define Optimizer/Loss function
- Learning rate scheduling
- Training
- Testing
- Run and Save model

# Set hyperparameters

```
74 # Training settings
75 parser = argparse.ArgumentParser(description='PyTorch MNIST Example')
76 parser.add_argument('--batch-size', type=int, default=64, metavar='N',
77                     help='input batch size for training (default: 64)')
78 parser.add_argument('--test-batch-size', type=int, default=1000, metavar='N',
79                     help='input batch size for testing (default: 1000)')
80 parser.add_argument('--epochs', type=int, default=14, metavar='N',
81                     help='number of epochs to train (default: 14)')
82 parser.add_argument('--lr', type=float, default=1.0, metavar='LR',
83                     help='learning rate (default: 1.0)')
84 parser.add_argument('--gamma', type=float, default=0.7, metavar='M',
85                     help='Learning rate step gamma (default: 0.7)')
86 parser.add_argument('--no-cuda', action='store_true', default=False,
87                     help='disables CUDA training')
88 parser.add_argument('--dry-run', action='store_true', default=False,
89                     help='quickly check a single pass')
90 parser.add_argument('--seed', type=int, default=1, metavar='S',
91                     help='random seed (default: 1)')
92 parser.add_argument('--log-interval', type=int, default=10, metavar='N',
93                     help='how many batches to wait before logging training status')
94 parser.add_argument('--save-model', action='store_true', default=False,
95                     help='For Saving the current Model')
96 args = parser.parse_args()
```

# Data Loader

- Pytorch offers data loaders for popular dataset

The following datasets are available:

## Datasets

- MNIST
- COCO
  - Captions
  - Detection
- LSUN
- ImageFolder
- Imagenet-12
- CIFAR
- STL10
- SVHN
- PhotoTour

# Data Loader

```
112 transform=transforms.Compose([
113     transforms.ToTensor(),
114     transforms.Normalize((0.1307,), (0.3081,))
115 ])
116 dataset1 = datasets.MNIST('../data', train=True, download=True,
117                             transform=transform)
118 dataset2 = datasets.MNIST('../data', train=False,
119                             transform=transform)
120 train_loader = torch.utils.data.DataLoader(dataset1, **train_kwargs)
121 test_loader = torch.utils.data.DataLoader(dataset2, **test_kwargs)
```

# Define Network

```
11 class Net(nn.Module):
12     def __init__(self):
13         super(Net, self).__init__()
14         self.conv1 = nn.Conv2d(1, 32, 3, 1)
15         self.conv2 = nn.Conv2d(32, 64, 3, 1)
16         self.dropout1 = nn.Dropout(0.25)
17         self.dropout2 = nn.Dropout(0.5)
18         self.fc1 = nn.Linear(9216, 128)
19         self.fc2 = nn.Linear(128, 10)
20
21     def forward(self, x):
22         x = self.conv1(x)
23         x = F.relu(x)
24         x = self.conv2(x)
25         x = F.relu(x)
26         x = F.max_pool2d(x, 2)
27         x = self.dropout1(x)
28         x = torch.flatten(x, 1)
29         x = self.fc1(x)
30         x = F.relu(x)
31         x = self.dropout2(x)
32         x = self.fc2(x)
33         output = F.log_softmax(x, dim=1)
34         return output
```

# Define Optimizer/Loss function

- Negative log likelihood loss

```
43         loss = F.nll_loss(output, target)
```

- Adadelta

```
124     optimizer = optim.Adadelta(model.parameters(), lr=args.lr)
```



# Learning rate scheduling

```
126 scheduler = StepLR(optimizer, step_size=1, gamma=args.gamma)
```

# Training

```
37 def train(args, model, device, train_loader, optimizer, epoch):
38     model.train()
39     for batch_idx, (data, target) in enumerate(train_loader):
40         data, target = data.to(device), target.to(device)
41         optimizer.zero_grad()
42         output = model(data)
43         loss = F.nll_loss(output, target)
44         loss.backward()
45         optimizer.step()
46         if batch_idx % args.log_interval == 0:
47             print('Train Epoch: {} [{}/{}] ({:.0f}%) \tLoss: {:.6f}'.format(
48                 epoch, batch_idx * len(data), len(train_loader.dataset),
49                 100. * batch_idx / len(train_loader), loss.item()))
50         if args.dry_run:
51             break
```

# Testing

```
54 def test(model, device, test_loader):
55     model.eval()
56     test_loss = 0
57     correct = 0
58     with torch.no_grad():
59         for data, target in test_loader:
60             data, target = data.to(device), target.to(device)
61             output = model(data)
62             test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch loss
63             pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
64             correct += pred.eq(target.view_as(pred)).sum().item()
65
66     test_loss /= len(test_loader.dataset)
67
68     print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%) \n'.format(
69         test_loss, correct, len(test_loader.dataset),
70         100. * correct / len(test_loader.dataset)))
```

# Run and Save model

```
127     for epoch in range(1, args.epochs + 1):
128         train(args, model, device, train_loader, optimizer, epoch)
129         test(model, device, test_loader)
130         scheduler.step()
131
132     if args.save_model:
133         torch.save(model.state_dict(), "mnist_cnn.pt")
```